constraints aligned with the survey coding scheme (e.g., one integer or one categorical option), enabling deterministic mapping from model outputs to the corresponding numeric variables [42]. In the *generic* setting (no cultural prompting), we ask the IVS survey question directly (with the same response-format constraints). In the *cultural* setting, we prepend the survey question with a lightweight country/territory prefix of the form "You are a citizen of X." to condition responses on the target cultural context.

Example prompting (generic vs. cultural).
A008 (Feeling of Happiness) Question from IVS

Generic (no cultural prompting):

```
Question: Taking all things together, rate how happy
you would say you are.
Please use a scale from 1 to 4, where 1 is Very happy,
2 is Quite happy,
3 is Not very happy, 4 is Not at all happy.
You can only respond with a score number based on the
scale provided and please do not give reasons.
Your score number:
```

Cultural prompting (country/territory-conditioned):

```
You are a citizen of United States of America (USA).
Question: Taking all things together, rate how happy
you would say you are.
Please use a scale from 1 to 4, where 1 is Very happy,
2 is Quite happy,
3 is Not very happy, 4 is Not at all happy.
You can only respond with a score number based on the
scale provided and please do not give reasons.
Your score number:
```

Following [42], we explicitly account for prompt-phrasing variance by evaluating a small set of semantically equivalent respondent-descriptor personas and averaging results to reduce sensitivity to wording choices; concretely, we vary the respondent descriptor using synonymous terms such as `average/typical`, `human being/person/individual`, and `world citizen` and average the resulting projected coordinates across variants .We use deterministic decoding (temperature = 0) and minimal output sanitation to extract the required option [42].

Let $\mathbf{x}_{m,c,v} \in \mathbb{R}^{10}$ be the coded response vector produced by model m under country/territory condition c and persona variant v, with $c = \varnothing$ denoting a non-national generic persona. We project each response vector into the IVS benchmark space by standardizing with IVS-derived moments and applying the IVS-fitted *rotated* two-dimensional scoring map:

$$\mathbf{z}_{m,c,v} = \left(\mathbf{x}_{m,c,v} - \boldsymbol{\mu}_{\text{raw}}^{\text{IVS}}\right) \oslash \boldsymbol{\sigma}_{\text{raw}}^{\text{IVS}}, \tag{3}$$

$$\mathbf{s}_{m,c,v} = \mathbf{W}_{\text{rot}}\, \mathbf{z}_{m,c,v}, \tag{4}$$

where $\boldsymbol{\mu}_{\text{raw}}^{\text{IVS}} \in \mathbb{R}^{10}$ and $\boldsymbol{\sigma}_{\text{raw}}^{\text{IVS}} \in \mathbb{R}^{10}$ are the (survey-weighted) IVS means and standard deviations for the ten indicators, $\oslash$ denotes elementwise division, and $\mathbf{W}_{\text{rot}} \in \mathbb{R}^{2\times 10}$ are the first-two-component *varimax-rotated* PCA scoring weights estimated from IVS data [42]. We then rescale $\mathbf{s}_{m,c,v}$ using Eqs. (1–2) to obtain $\boldsymbol{\pi}_{m,c,v} \in \mathbb{R}^2$ and average across persona variants:

$$\boldsymbol{\mu}_{m,c} = \frac{1}{|V|} \sum_{v \in V} \boldsymbol{\pi}_{m,c,v}. \tag{5}$$

We overlay $\boldsymbol{\mu}_{m,\varnothing}$ (generic, non-national prompting) for each open-weight model on the IVS cultural map to visualize how open-weight models cluster relative to country/territory reference locations (Figure 1).

3.3 Country-level cultural distance under three prompting regimes

Having embedded both IVS respondents and model responses in a common space, we quantify alignment using country-level Euclidean distance in the IVS benchmark coordinates as described in [42]. Let $\boldsymbol{\nu}_c^{\text{IVS}} \in \mathbb{R}^2$ denote the IVS reference coordinate for country/territory c (computed from human respondents). All distances are measured to the same human reference point $\boldsymbol{\nu}_c^{\text{IVS}}$.

We compare three prompting regimes for cultural conditioning:

(1) No culture conditioning. We compute a single generic model point $\boldsymbol{\mu}_{m,\varnothing}$ and compare it to every country benchmark:

$$d_{\text{nocond}}(m,c) = \left\|\boldsymbol{\mu}_{m,\varnothing} - \boldsymbol{\nu}_c^{\text{IVS}}\right\|_2. \tag{6}$$

(2) Manual culture prompting. We condition the model on country/territory identity c via a fixed cultural prefix while keeping the survey questions and answer constraints unchanged [42]. This yields $\boldsymbol{\mu}_{m,c}^{\text{man}}$ and distance

$$d_{\text{man}}(m,c) = \left\|\boldsymbol{\mu}_{m,c}^{\text{man}} - \boldsymbol{\nu}_c^{\text{IVS}}\right\|_2. \tag{7}$$

(3) Culture prompt programming (*DSPy*). We replace the fixed manual prefix with a compiled prompt program produced by DSPy [24], as described in Section 3.4, yielding $\boldsymbol{\mu}_{m,c}^{\text{DSPy}}$ and distance

$$d_{\text{DSPy}}(m,c) = \left\|\boldsymbol{\mu}_{m,c}^{\text{DSPy}} - \boldsymbol{\nu}_c^{\text{IVS}}\right\|_2. \tag{8}$$

To summarize changes relative to the no-conditioning baseline, we report paired country-level improvements

$$\Delta d_{\text{man}}(m,c) = d_{\text{man}}(m,c) - d_{\text{nocond}}(m,c), \tag{9}$$

$$\Delta d_{\text{DSPy}}(m,c) = d_{\text{DSPy}}(m,c) - d_{\text{nocond}}(m,c), \tag{10}$$

and the fraction of countries with $\Delta d_{\text{man}}(m,c) < 0$ and $\Delta d_{\text{DSPy}}(m,c) < 0$. We visualize the distances in Figure 2.

3.4 Culture prompt programming with DSPy

We use *DSPy* for prompt programming because it provides a way to treat cultural conditioning as an optimization problem: instead of committing to a single hand-written persona template, we allow the culture-conditioning instruction to be automatically tuned to reduce country-level cultural distance [24]. Concretely, we parameterize the culture-conditioning instruction by a discrete prompt parameter θ (a text instruction) that is inserted upstream of the fixed IVS question block; for a given (m, c, θ), the resulting program elicits the ten IVS responses under the same response-format constraints used throughout and projects them into the IVS space to obtain $\boldsymbol{\mu}_{m,c}(\theta) \in \mathbb{R}^2$ (Sections 3.2–3.3) [42].

We train the compiler to reduce misalignment by directly optimizing country-level cultural distance to IVS benchmarks. Let $\boldsymbol{\nu}_c^{\text{IVS}} \in \mathbb{R}^2$ be the IVS benchmark coordinate for country/territory